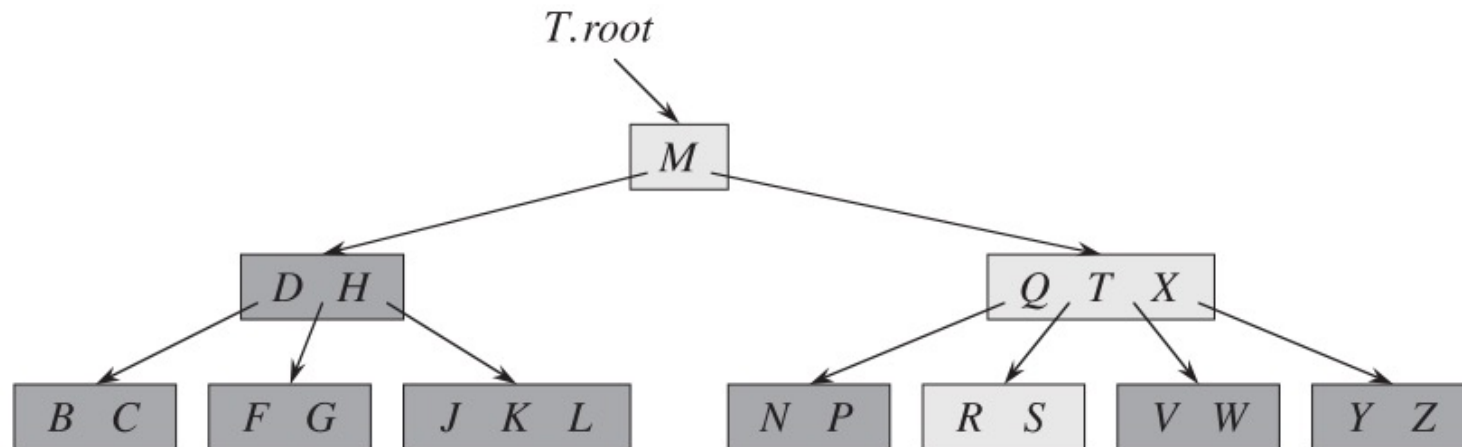


B-Trees

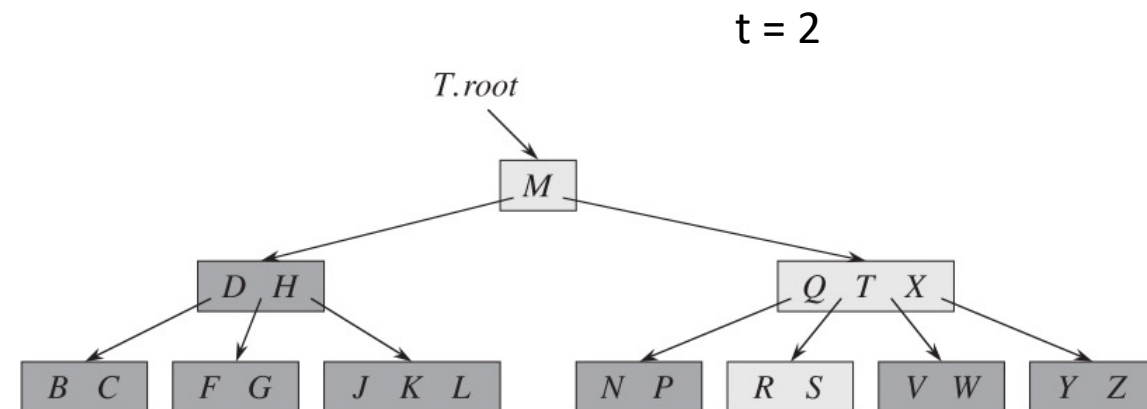
What is a B-Tree?

- Balanced search tree
- Each B-tree node can have many children
 - Branching factor: a few to thousands
- Height: $O(\lg n)$




Properties of B-trees

- Keys are stored in non-decreasing order in a node
- All leaves have the same depth – tree height h
- For each non-leaf nodes with i elements contains $i+1$ pointers
- **t : minimum degree**, integer, $t \geq 2$
 - Every non-root node must have $\geq t-1$ keys
 - Every non-root internal node have $\geq t$ children
 - Every node contains $\leq 2t-1$ keys
 - Every internal nodes contain $\leq 2t$ children



Properties of B-trees

- **t: minimum degree**, integer, $t \geq 2$
 - Every non-root node must have $\geq t-1$ keys
 - Every non-root internal node have $\geq t$ children
 - Every node contains $\leq 2t-1$ keys
 - Every internal nodes contain $\leq 2t$ children
- Can be summarized in 3 main rules:
 - Every none root node must have $\lceil t/2 \rceil$ (Ceiling of $t/2$) children
 - Root must have min 2 children
 - All leaf nodes must be at the same level.



B-tree search

Search tree rooted at node x for key k

B-TREE-SEARCH(x, k)

```

1   $i = 1$ 
2  while  $i \leq x.n$  and  $k > x.key_i$ 
3       $i = i + 1$ 
4  if  $i \leq x.n$  and  $k == x.key_i$ 
5      return  $(x, i)$ 
6  elseif  $x.leaf$ 
7      return NIL
8  else DISK-READ( $x.c_i$ )
9      return B-TREE-SEARCH( $x.c_i, k$ )
    
```

} find smallest index i in the tree where $key[i] \geq k$
 } if $k == key[i]$, found the key and return the results
 } if $k \neq key[i]$ and x is a leaf, could not find k
 } if $k \neq key[i]$ and x is not a leaf, recurse and search the subtree of x

pointer

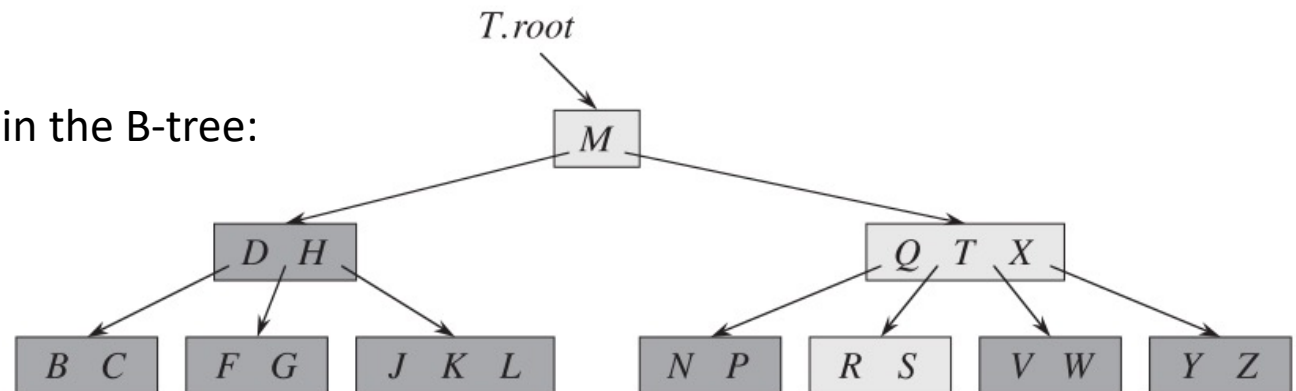
height

complexity: $O(t h) = O(t \log_t n)$

minimum
degree

num of
keys

Search for R in the B-tree:



Create an empty B-tree

B-TREE-CREATE(T)

1 $x = \text{ALLOCATE-NODE}()$

2 $x.\textit{leaf} = \text{TRUE}$

3 $x.n = 0$

4 $\text{DISK-WRITE}(x)$

5 $T.\textit{root} = x$

Insert a key into B-trees – Splitting a node

Before inserting, if the node is full with $2t-1$ keys, need to split the node first

B-TREE-SPLIT-CHILD(x, i) ← split the full node $x.c_i$, which is y , where x is a nonfull internal node



```
1  z = ALLOCATE-NODE()
2  y = x.ci
3  z.leaf = y.leaf
4  z.n = t - 1
5  for j = 1 to t - 1
6      z.keyj = y.keyj+t
7  if not y.leaf
8      for j = 1 to t
9          z.cj = y.cj+t
10 y.n = t - 1
11 for j = x.n + 1 downto i + 1
12     x.cj+1 = x.cj
13 x.ci+1 = z
14 for j = x.n downto i
15     x.keyj+1 = x.keyj
16 x.keyi = y.keyt
17 x.n = x.n + 1
18 DISK-WRITE(y)
19 DISK-WRITE(z)
20 DISK-WRITE(x)
```

y has $2t-1$ elements.

This creates node z and allocate the largest $t-1$ elements in y into z , and adjust key count for y

Assign z as a child of x , move the median key from y to x , and adjust key count of x

complexity: $\Theta(t)$



Insert a key into B-trees – nonfull node


B-TREE-INSERT-NONFULL(x, k) \longrightarrow insert a key k into a non-full node x

```
1   $i = x.n$ 
2  if  $x.leaf$ 
3      while  $i \geq 1$  and  $k < x.key_i$ 
4           $x.key_{i+1} = x.key_i$ 
5           $i = i - 1$ 
6       $x.key_{i+1} = k$ 
7       $x.n = x.n + 1$ 
8      DISK-WRITE( $x$ )
9  else while  $i \geq 1$  and  $k < x.key_i$ 
10      $i = i - 1$ 
11      $i = i + 1$ 
12     DISK-READ( $x.c_i$ )
13     if  $x.c_i.n == 2t - 1$ 
14         B-TREE-SPLIT-CHILD( $x, i$ )
15         if  $k > x.key_i$ 
16              $i = i + 1$ 
17     B-TREE-INSERT-NONFULL( $x.c_i, k$ )
```

if x is a leaf node, then insert k into x

if x is not a leaf node, and if $x.c_i$ is full, then split $x.c_i$, and find which child node to insert k into

recursively insert k into that nonfull child node



Insert a key into B-trees – general case

B-TREE-INSERT(T, k) \longrightarrow insert new key k into B-tree T

```


1   $r = T.root$ 
2  if  $r.n == 2t - 1$ 
3       $s = \text{ALLOCATE-NODE}()$ 
4       $T.root = s$ 
5       $s.leaf = \text{FALSE}$ 
6       $s.n = 0$ 
7       $s.c_1 = r$ 
8      B-TREE-SPLIT-CHILD( $s, 1$ )
9      B-TREE-INSERT-NONFULL( $s, k$ )
10 else B-TREE-INSERT-NONFULL( $r, k$ )
    
```

if the root node r is full, split the root

insert k into the new nonfull root by calling B-tree-Insert-Nonfull


if root is nonfull, insert k into the root by calling B-tree-Insert-Nonfull

(c) Q inserted







insert L into the B-tree \longrightarrow

(d) L inserted



Delete from B-Tree

$t = 3$



initial case

Case 1: simple deletion from leaf node


Case 2a: If child y preceding k has at least t keys, float predecessor k' to replace k , do it recursively

Case 2b: if y has less than t keys, find z that follows k , if z has at least t keys, float the successor of k' of k to replace k , do it recursively


Case 2c: otherwise, merge k and all of z into y and recursively delete k from y

Delete from B-tree continued


(e) *D* deleted: case 3b



(e') tree shrinks in height



(f) *B* deleted: case 3a



Case 3a: if *x.ci* has only $t-1$ keys but immediate sibling has at least t keys, move a key from *x* into *ci* and move a key from *x.ci* sibling up into *x*, and then delete *k*, recurse on child of *x*

Case 3b: if *x.ci* and *x.ci*'s immediate siblings have $t-1$ keys, merge *x.ci* with one sibling and move a key from *x* down to the median of the new node, do it recursively

Reading

- CLRS Book.

Introduction to Algorithms, 3rd Edition, Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, and Clifford Stein

B-Tree Visualization

<https://www.cs.usfca.edu/~galles/visualization/BTree.html>